

project_5_starter

March 20, 2023

1 Project 5: NLP on Financial Statements

1.1 Instructions

Each problem consists of a function to implement and instructions on how to implement the function. The parts of the function that need to be implemented are marked with a `# TODO` comment. After implementing the function, run the cell to test it against the unit tests we've provided. For each problem, we provide one or more unit tests from our `project_tests` package. These unit tests won't tell you if your answer is correct, but will warn you of any major errors. Your code will be checked for the correct solution when you submit it to Udacity.

1.2 Packages

When you implement the functions, you'll only need to use the packages you've used in the classroom, like [Pandas](#) and [Numpy](#). These packages will be imported for you. We recommend you don't add any import statements, otherwise the grader might not be able to run your code.

The other packages that we're importing are `project_helper` and `project_tests`. These are custom packages built to help you solve the problems. The `project_helper` module contains utility functions and graph functions. The `project_tests` contains the unit tests for all the problems.

1.2.1 Install Packages

```
In [2]: import sys
```

```
!{sys.executable} -m pip install -r requirements.txt --user
```

```
Collecting alphalens==0.3.2 (from -r requirements.txt (line 1))
```

```
  Downloading https://files.pythonhosted.org/packages/a5/dc/2f9cd107d0d4cf6223d37d81ddfbdbbf0d70
```

```
  100% || 18.9MB 1.8MB/s eta 0:00:01
```

```
Collecting nltk==3.3.0 (from -r requirements.txt (line 2))
```

```
  Downloading https://files.pythonhosted.org/packages/50/09/3b1755d528ad9156ee7243d52aa5cd2b809e
```

```
  100% || 1.4MB 15.4MB/s ta 0:00:01
```

```
Collecting numpy==1.13.3 (from -r requirements.txt (line 3))
```

```
  Downloading https://files.pythonhosted.org/packages/57/a7/e3e6bd9d595125e1abbe162e323fd2d06f6f
```

```
  100% || 17.0MB 2.4MB/s eta 0:00:01
```

```
Collecting ratelimit==2.2.0 (from -r requirements.txt (line 4))
```

```
  Downloading https://files.pythonhosted.org/packages/b5/73/956d739706da2f74891ba46391381ce7e680
```

```
Requirement already satisfied: requests==2.18.4 in /opt/conda/lib/python3.6/site-packages (from
```

```
Requirement already satisfied: scikit-learn==0.19.1 in /opt/conda/lib/python3.6/site-packages (f
```

```

Requirement already satisfied: six==1.11.0 in /opt/conda/lib/python3.6/site-packages (from -r re
Collecting tqdm==4.19.5 (from -r requirements.txt (line 8))
  Downloading https://files.pythonhosted.org/packages/71/3c/341b4fa23cb3abc335207dba057c790f3bb3
100% || 61kB 17.7MB/s ta 0:00:01
Requirement already satisfied: matplotlib>=1.4.0 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: pandas>=0.18.0 in /opt/conda/lib/python3.6/site-packages (from al
Requirement already satisfied: scipy>=0.14.0 in /opt/conda/lib/python3.6/site-packages (from alp
Requirement already satisfied: seaborn>=0.6.0 in /opt/conda/lib/python3.6/site-packages (from al
Requirement already satisfied: statsmodels>=0.6.1 in /opt/conda/lib/python3.6/site-packages (fro
Requirement already satisfied: IPython>=3.2.3 in /opt/conda/lib/python3.6/site-packages (from al
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /opt/conda/lib/python3.6/site-packages (
Requirement already satisfied: idna<2.7,>=2.5 in /opt/conda/lib/python3.6/site-packages (from re
Requirement already satisfied: urllib3<1.23,>=1.21.1 in /opt/conda/lib/python3.6/site-packages (
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.6/site-packages (fro
Requirement already satisfied: python-dateutil>=2.0 in /opt/conda/lib/python3.6/site-packages (f
Requirement already satisfied: pytz in /opt/conda/lib/python3.6/site-packages (from matplotlib>=
Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.6/site-packages/cycler-0.1
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in /opt/conda/lib/python
Requirement already satisfied: traitlets>=4.2 in /opt/conda/lib/python3.6/site-packages (from IP
Requirement already satisfied: decorator in /opt/conda/lib/python3.6/site-packages (from IPython
Requirement already satisfied: simplegeneric>0.8 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: pygments in /opt/conda/lib/python3.6/site-packages (from IPython>
Requirement already satisfied: jedi>=0.10 in /opt/conda/lib/python3.6/site-packages (from IPytho
Requirement already satisfied: pickleshare in /opt/conda/lib/python3.6/site-packages (from IPyth
Requirement already satisfied: setuptools>=18.5 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: backcall in /opt/conda/lib/python3.6/site-packages (from IPython>
Requirement already satisfied: pexpect; sys_platform != "win32" in /opt/conda/lib/python3.6/site
Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.15 in /opt/conda/lib/python3.6/site-pa
Requirement already satisfied: ipython-genutils in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: ptyprocess>=0.5 in /opt/conda/lib/python3.6/site-packages (from p
Requirement already satisfied: wcwidth in /opt/conda/lib/python3.6/site-packages (from prompt-to
Building wheels for collected packages: alphasens, nltk, ratelimit
  Running setup.py bdist_wheel for alphasens ... done
  Stored in directory: /home/student/.cache/pip/wheels/77/1e/9a/223b4c94d7f564f25d94b48ca5b9c53e
  Running setup.py bdist_wheel for nltk ... done
  Stored in directory: /home/student/.cache/pip/wheels/d1/ab/40/3bceea46922767e42986aef7606a6005
  Running setup.py bdist_wheel for ratelimit ... done
  Stored in directory: /home/student/.cache/pip/wheels/a6/2a/13/3c6e42757ca0b6873a60e0697d30f7dd
Successfully built alphasens nltk ratelimit
tensorflow 1.3.0 requires tensorflow-tensorboard<0.2.0,>=0.1.0, which is not installed.
moviepy 0.2.3.2 has requirement tqdm==4.11.2, but you'll have tqdm 4.19.5 which is incompatible.
Installing collected packages: numpy, alphasens, nltk, ratelimit, tqdm
  The script tqdm is installed in '/home/student/.local/bin' which is not on PATH. Consider add
Successfully installed alphasens-0.3.2 nltk-3.3 numpy-1.13.3 ratelimit-2.2.0 tqdm-4.19.5

```

1.2.2 Load Packages

```
In [2]: import nltk
import numpy as np
import pandas as pd
import pickle
import pprint
import csv
import ast

import project_helper
import project_tests

from tqdm import tqdm
from bs4 import BeautifulSoup
from collections import defaultdict
```

1.2.3 Download NLP Corpora

You'll need two corpora to run this project: the stopwords corpus for removing stopwords and wordnet for lemmatizing.

```
In [3]: nltk.download('stopwords')
nltk.download('wordnet')
```

[nltk_data] Downloading package stopwords to
[nltk_data] /home/student/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package wordnet to /home/student/nltk_data...

```
Out[3]: True
```

1.3 Get 10ks

We'll be running NLP analysis on 10-k documents. To do that, we first need to download the documents. For this project, we'll download 10-ks for a few companies. To lookup documents for these companies, we'll use their CIK. If you would like to run this against other stocks, we've provided the dict `additional_cik` for more stocks. However, the more stocks you try, the long it will take to run.

```
In [4]: cik_lookup = {
    'AMZN': '0001018724',
    'BMY': '0000014272',
    'CNP': '0001130310',
    'CVX': '0000093410',
    'FL': '0000850209',
    'FRT': '0000034903',
    'HON': '0000773840'}
```

```

additional_cik = {
    'AEP': '0000004904',
    'AXP': '0000004962',
    'BA': '0000012927',
    'BK': '0001390777',
    'CAT': '0000018230',
    'DE': '0000315189',
    'DIS': '0001001039',
    'DTE': '0000936340',
    'ED': '0001047862',
    'EMR': '0000032604',
    'ETN': '0001551182',
    'GE': '0000040545',
    'IBM': '0000051143',
    'IP': '0000051434',
    'JNJ': '0000200406',
    'KO': '0000021344',
    'LLY': '0000059478',
    'MCD': '0000063908',
    'MO': '0000764180',
    'MRK': '0000310158',
    'MRO': '0000101778',
    'PCG': '0001004980',
    'PEP': '0000077476',
    'PFE': '0000078003',
    'PG': '0000080424',
    'PNR': '0000077360',
    'SYY': '0000096021',
    'TXN': '0000097476',
    'UTX': '0000101829',
    'WFC': '0000072971',
    'WMT': '0000104169',
    'WY': '0000106535',
    'XOM': '0000034088'}

```

1.3.1 Get list of 10-ks

The SEC has a limit on the number of calls you can make to the website per second. In order to avoid hitting that limit, we've created the SecAPI class. This will cache data from the SEC and prevent you from going over the limit.

```
In [5]: sec_api = project_helper.SecAPI()
```

With the class constructed, let's pull a list of filled 10-ks from the SEC for each company.

1.3.2 Upload Stock RSS URL from csv file

```
In [6]: sec_data = {}
```

```
    with open('sec_data.csv', 'r') as f:
        reader = csv.reader(f)
        for row in reader:
            ticker, rss_url = row
            sec_data[ticker] = ast.literal_eval(rss_url)
```

```
In [7]: example_ticker = 'AMZN'
        pprint.pprint(sec_data[example_ticker][:5])
```

```
[('https://www.sec.gov/Archives/edgar/data/1018724/000101872417000011/0001018724-17-000011-index
  '10-K',
  '2017-02-10'),
 ('https://www.sec.gov/Archives/edgar/data/1018724/000101872416000172/0001018724-16-000172-index
  '10-K',
  '2016-01-29'),
 ('https://www.sec.gov/Archives/edgar/data/1018724/000101872415000006/0001018724-15-000006-index
  '10-K',
  '2015-01-30'),
 ('https://www.sec.gov/Archives/edgar/data/1018724/000101872414000006/0001018724-14-000006-index
  '10-K',
  '2014-01-31'),
 ('https://www.sec.gov/Archives/edgar/data/1018724/000119312513028520/0001193125-13-028520-index
  '10-K',
  '2013-01-30')]
```

1.3.3 Download 10-ks

As you see, this is a list of urls. These urls point to a file that contains metadata related to each filling. Since we don't care about the metadata, we'll pull the filling by replacing the url with the filling url.

1.3.4 get 10-k data from a csv file

```
In [8]: pd_fillings_by_ticker = {}
        colnames=['ticker', 'file_date', '10-k']
```

```
        pd_fillings_by_ticker = pd.read_csv('raw_fillings_by_ticker.csv',
                                              header=None,
                                              names=colnames)
```

```
In [9]: pd_fillings_by_ticker.head()
```

```
Out[9]:   ticker  file_date                                     10-k
0    AMZN  2017-02-10  <SEC-DOCUMENT>0001018724-17-000011.txt : 20170...
```

```

1  AMZN  2016-01-29  <SEC-DOCUMENT>0001018724-16-000172.txt : 20160...
2  AMZN  2015-01-30  <SEC-DOCUMENT>0001018724-15-000006.txt : 20150...
3  AMZN  2014-01-31  <SEC-DOCUMENT>0001018724-14-000006.txt : 20140...
4  AMZN  2013-01-30  <SEC-DOCUMENT>0001193125-13-028520.txt : 20130...

```

```
In [10]: raw_fillings_by_ticker = defaultdict(dict)
```

```

    for i, row in pd_fillings_by_ticker.iterrows():
        raw_fillings_by_ticker[row.ticker][row.file_date] = row['10-k']

```

```
In [11]: print('Example Document:\n\n{}'.format(next(iter(raw_fillings_by_ticker[example_ticker])))
```

Example Document:

```

<SEC-DOCUMENT>0001018724-17-000011.txt : 20170210
<SEC-HEADER>0001018724-17-000011.hdr.sgml : 20170210
<ACCEPTANCE-DATETIME>20170209175636
ACCESSION NUMBER:          0001018724-17-000011
CONFORMED SUBMISSION TYPE: 10-K
PUBLIC DOCUMENT COUNT:     92
CONFORMED PERIOD OF REPORT: 20161231
FILED AS OF DATE:          20170210
DATE AS OF CHANGE:         20170209

```

FILER:

COMPANY DATA:

```

    COMPANY CONFORMED NAME:          AMAZON COM INC
    CENTRAL INDEX KEY:                0001018724
    STANDARD INDUSTRIAL CLASSIFICATION: RETAIL-CATALOG & MAIL-ORDER HOUSES [5
    IRS NUMBER:                      911646860
    STATE OF INCORPORATION:          DE
    FISCAL YEAR END:                 1231

```

FILING VALUES:

```

    FORM TYPE:          10-K
    SEC ACT:            1934 Act
    SEC FILE NUMBER:    000-22513
    FILM NUMBER:        17588807

```

BUSINESS ADDRESS:

```

    STREET 1:          410 TERRY AVENUE NORTH
    CITY:              SEATTLE
    STATE:             WA
    ZIP:               98109
    BUSINESS PHONE:    2062661000

```

MAIL ADDRESS:

```

        STREET 1:          410 TERRY AVENUE NORTH
        CITY:              SEATTLE
        STATE:             WA
        ZIP:               98109
</SEC-HEADER>
<DOCUMENT>
<TYPE>10-K
<SEQUENCE>1
<FILENAME...

```

1.3.5 Get Documents

With theses fillings downloaded, we want to break them into their associated documents. These documents are sectioned off in the fillings with the tags <DOCUMENT> for the start of each document and </DOCUMENT> for the end of each document. There's no overlap with these documents, so each </DOCUMENT> tag should come after the <DOCUMENT> with no <DOCUMENT> tag in between.

Implement `get_documents` to return a list of these documents from a filling. Make sure not to include the tag in the returned document text.

```
In [12]: import re
```

```

def get_documents(text):
    """
    Extract the documents from the text

    Parameters
    -----
    text : str
        The text with the document strings inside

    Returns
    -----
    extracted_docs : list of str
        The document strings found in `text`
    """

    # TODO: Implement
    pattern = r"<DOCUMENT>(.*?)</DOCUMENT>"
    extracted_docs = re.findall(pattern, text, re.DOTALL)

    return extracted_docs

project_tests.test_get_documents(get_documents)

```

Tests Passed

With the `get_documents` function implemented, let's extract all the documents.

```
In [13]: filling_documents_by_ticker = {}

for ticker, raw_fillings in raw_fillings_by_ticker.items():
    filling_documents_by_ticker[ticker] = {}
    for file_date, filling in tqdm(raw_fillings.items(), desc='Getting Documents from {}'):
        filling_documents_by_ticker[ticker][file_date] = get_documents(filling)

print('\n\n'.join([
    'Document {} Filed on {}: \n{}'.format(doc_i, file_date, doc[:200])
    for file_date, docs in filling_documents_by_ticker[example_ticker].items()
    for doc_i, doc in enumerate(docs)[:3]))
```

```
Getting Documents from AMZN Fillings: 100%|| 17/17 [00:01<00:00, 9.30filling/s]
Getting Documents from BMJ Fillings: 100%|| 23/23 [00:04<00:00, 5.03filling/s]
Getting Documents from CNP Fillings: 100%|| 15/15 [00:03<00:00, 4.33filling/s]
Getting Documents from CVX Fillings: 100%|| 21/21 [00:04<00:00, 4.43filling/s]
Getting Documents from FL Fillings: 100%|| 16/16 [00:02<00:00, 7.13filling/s]
Getting Documents from FRT Fillings: 100%|| 19/19 [00:02<00:00, 6.63filling/s]
Getting Documents from HON Fillings: 100%|| 20/20 [00:03<00:00, 6.02filling/s]
```

Document 0 Filed on 2017-02-10:

```
<TYPE>10-K
<SEQUENCE>1
<FILENAME>amzn-20161231x10k.htm
<DESCRIPTION>FORM 10-K
<TEXT>
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http://www.w3.org/TR/html4/loose
<html>

    <he...
```

Document 1 Filed on 2017-02-10:

```
<TYPE>EX-12.1
<SEQUENCE>2
<FILENAME>amzn-20161231xex121.htm
<DESCRIPTION>COMPUTATION OF RATIO OF EARNINGS TO FIXED CHARGES
<TEXT>
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http:...
```

Document 2 Filed on 2017-02-10:

```
<TYPE>EX-21.1
<SEQUENCE>3
<FILENAME>amzn-20161231xex211.htm
```



```
<DESCRIPTION>LIST OF SIGNIFICANT SUBSIDIARIES
<TEXT>
<!DOCTYPE html PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN" "http://www.w3.org/TR/h...
```

1.3.6 Get Document Types

Now that we have all the documents, we want to find the 10-k form in this 10-k filing. Implement the `get_document_type` function to return the type of document given. The document type is located on a line with the `<TYPE>` tag. For example, a form of type "TEST" would have the line `<TYPE>TEST`. Make sure to return the type as lowercase, so this example would be returned as "test".

```
In [14]: def get_document_type(doc):
        """
        Return the document type lowercased

        Parameters
        -----
        doc : str
            The document string

        Returns
        -----
        doc_type : str
            The document type lowercased
        """

        # TODO: Implement
        pattern = r"<TYPE>[^\n]+"
        match = re.findall(pattern, doc)
        doc_type = str.lower(match[0][len("<TYPE>"):])

        return doc_type

project_tests.test_get_document_type(get_document_type)
```

Tests Passed

With the `get_document_type` function, we'll filter out all non 10-k documents.

```
In [15]: ten_ks_by_ticker = {}

        for ticker, filling_documents in filling_documents_by_ticker.items():
            ten_ks_by_ticker[ticker] = []
```

```

        for file_date, documents in filling_documents.items():
            for document in documents:
                if get_document_type(document) == '10-k':
                    ten_ks_by_ticker[ticker].append({
                        'cik': cik_lookup[ticker],
                        'file': document,
                        'file_date': file_date})

    project_helper.print_ten_k_data(ten_ks_by_ticker[example_ticker][:5], ['cik', 'file', 'file_date'])

[
  {
    cik: '0001018724'
    file: '\n<TYPE>10-K\n<SEQUENCE>1\n<FILENAME>amzn-2016123...
    file_date: '2017-02-10'},
  {
    cik: '0001018724'
    file: '\n<TYPE>10-K\n<SEQUENCE>1\n<FILENAME>amzn-2015123...
    file_date: '2016-01-29'},
  {
    cik: '0001018724'
    file: '\n<TYPE>10-K\n<SEQUENCE>1\n<FILENAME>amzn-2014123...
    file_date: '2015-01-30'},
  {
    cik: '0001018724'
    file: '\n<TYPE>10-K\n<SEQUENCE>1\n<FILENAME>amzn-2013123...
    file_date: '2014-01-31'},
  {
    cik: '0001018724'
    file: '\n<TYPE>10-K\n<SEQUENCE>1\n<FILENAME>d445434d10k...
    file_date: '2013-01-30'},
]

```

1.4 Preprocess the Data

1.4.1 Clean Up

As you can see, the text for the documents are very messy. To clean this up, we'll remove the html and lowercase all the text.

```

In [16]: def remove_html_tags(text):
            text = BeautifulSoup(text, 'html.parser').get_text()

            return text

        def clean_text(text):

```

```

text = text.lower()
text = remove_html_tags(text)

return text

```

Using the `clean_text` function, we'll clean up all the documents.

```

In [17]: for ticker, ten_ks in ten_ks_by_ticker.items():
        for ten_k in tqdm(ten_ks, desc='Cleaning {} 10-Ks'.format(ticker), unit='10-K'):
            ten_k['file_clean'] = clean_text(ten_k['file'])

project_helper.print_ten_k_data(ten_ks_by_ticker[example_ticker][:5], ['file_clean'])

Cleaning AMZN 10-Ks: 100%|| 17/17 [00:28<00:00, 1.65s/10-K]
Cleaning BMY 10-Ks: 100%|| 23/23 [01:04<00:00, 2.80s/10-K]
Cleaning CNP 10-Ks: 100%|| 15/15 [00:48<00:00, 3.21s/10-K]
Cleaning CVX 10-Ks: 100%|| 21/21 [01:33<00:00, 4.48s/10-K]
Cleaning FL 10-Ks: 100%|| 16/16 [00:21<00:00, 1.32s/10-K]
Cleaning FRT 10-Ks: 100%|| 19/19 [00:45<00:00, 2.38s/10-K]
Cleaning HON 10-Ks: 100%|| 19/19 [00:45<00:00, 2.39s/10-K]

[
  {
    file_clean: '\n10-k\n1\namzn-20161231x10k.htm\nform 10-k\n\n\n...',
  },
  {
    file_clean: '\n10-k\n1\namzn-20151231x10k.htm\nform 10-k\n\n\n...',
  },
  {
    file_clean: '\n10-k\n1\namzn-20141231x10k.htm\nform 10-k\n\n\n...',
  },
  {
    file_clean: '\n10-k\n1\namzn-20131231x10k.htm\nform 10-k\n\n\n...',
  },
  {
    file_clean: '\n10-k\n1\nd445434d10k.htm\nform 10-k\n\n\nform 1...',
  },
]

```

1.4.2 Lemmatize

With the text cleaned up, it's time to distill the verbs down. Implement the `lemmatize_words` function to lemmatize verbs in the list of words provided.

```

In [18]: from nltk.stem import WordNetLemmatizer
        from nltk.corpus import wordnet

        def lemmatize_words(words):

```

```

"""
Lemmatize words

Parameters
-----
words : list of str
    List of words

Returns
-----
lemmatized_words : list of str
    List of lemmatized words
"""

# TODO: Implement
lemmatized_words = [WordNetLemmatizer().lemmatize(w, pos='v') for w in words]

return lemmatized_words

```

```
project_tests.test_lemmatize_words(lemmatize_words)
```

Tests Passed

With the `lemmatize_words` function implemented, let's lemmatize all the data.

```

In [ ]: word_pattern = re.compile('\w+')

for ticker, ten_ks in ten_ks_by_ticker.items():
    for ten_k in tqdm(ten_ks, desc='Lemmatize {} 10-Ks'.format(ticker), unit='10-K'):
        ten_k['file_lemma'] = lemmatize_words(word_pattern.findall(ten_k['file_clean']))

```

```
project_helper.print_ten_k_data(ten_ks_by_ticker[example_ticker][:5], ['file_lemma'])
```

```

Lemmatize AMZN 10-Ks: 100%|| 17/17 [00:04<00:00, 4.0610-K/s]
Lemmatize BMY 10-Ks: 100%|| 23/23 [00:09<00:00, 2.5110-K/s]
Lemmatize CNP 10-Ks: 100%|| 15/15 [00:07<00:00, 2.0210-K/s]
Lemmatize CVX 10-Ks: 52%| | 11/21 [00:04<00:04, 2.2410-K/s]

```

1.4.3 Remove Stopwords

```

In [ ]: from nltk.corpus import stopwords

lemma_english_stopwords = lemmatize_words(stopwords.words('english'))

for ticker, ten_ks in ten_ks_by_ticker.items():

```

```

for ten_k in tqdm(ten_ks, desc='Remove Stop Words for {} 10-Ks'.format(ticker), unit=
    ten_k['file_lemma'] = [word for word in ten_k['file_lemma'] if word not in lemma

print('Stop Words Removed')

```

1.5 Analysis on 10ks

1.5.1 Loughran McDonald Sentiment Word Lists

We'll be using the Loughran and McDonald sentiment word lists. These word lists cover the following sentiment: - Negative - Positive - Uncertainty - Litigious - Constraining - Superfluous - Modal

This will allow us to do the sentiment analysis on the 10-ks. Let's first load these word lists. We'll be looking into a few of these sentiments.

```

In [ ]: import os

sentiments = ['negative', 'positive', 'uncertainty', 'litigious', 'constraining', 'inter

sentiment_df = pd.read_csv(os.path.join('..', '..', 'data', 'project_5_loughran_mcdonald
sentiment_df.columns = [column.lower() for column in sentiment_df.columns] # Lowercase t

# Remove unused information
sentiment_df = sentiment_df[sentiments + ['word']]
sentiment_df[sentiments] = sentiment_df[sentiments].astype(bool)
sentiment_df = sentiment_df[(sentiment_df[sentiments]).any(1)]

# Apply the same preprocessing to these words as the 10-k words
sentiment_df['word'] = lemmatize_words(sentiment_df['word'].str.lower())
sentiment_df = sentiment_df.drop_duplicates('word')

sentiment_df.head()

```

1.5.2 Bag of Words

using the sentiment word lists, let's generate sentiment bag of words from the 10-k documents. Implement `get_bag_of_words` to generate a bag of words that counts the number of sentiment words in each doc. You can ignore words that are not in `sentiment_words`.

```

In [ ]: from collections import defaultdict, Counter
        from sklearn.feature_extraction.text import CountVectorizer

def get_bag_of_words(sentiment_words, docs):
    """
    Generate a bag of words from documents for a certain sentiment

```

```

Parameters
-----
sentiment_words: Pandas Series
    Words that signify a certain sentiment
docs : list of str
    List of documents used to generate bag of words

Returns
-----
bag_of_words : 2-d Numpy Nddarray of int
    Bag of words sentiment for each document
    The first dimension is the document.
    The second dimension is the word.
"""

# TODO: Implement
cnt_vec = CountVectorizer(vocabulary=sentiment_words)
bag_of_words = cnt_vec.fit_transform(docs).toarray()

return bag_of_words

```

```
project_tests.test_get_bag_of_words(get_bag_of_words)
```

Using the `get_bag_of_words` function, we'll generate a bag of words for all the documents.

```

In [ ]: sentiment_bow_ten_ks = {}

for ticker, ten_ks in ten_ks_by_ticker.items():
    lemma_docs = [' '.join(ten_k['file_lemma']) for ten_k in ten_ks]

    sentiment_bow_ten_ks[ticker] = {
        sentiment: get_bag_of_words(sentiment_df[sentiment_df[sentiment]]['word'], lemma_docs)
        for sentiment in sentiments}

project_helper.print_ten_k_data([sentiment_bow_ten_ks[example_ticker]], sentiments)

```

1.5.3 Jaccard Similarity

Using the bag of words, let's calculate the jaccard similarity on the bag of words and plot it over time. Implement `get_jaccard_similarity` to return the jaccard similarities between each tick in time. Since the input, `bag_of_words_matrix`, is a bag of words for each time period in order, you just need to compute the jaccard similarities for each neighboring bag of words. Make sure to turn the bag of words into a boolean array when calculating the jaccard similarity.

```
In [ ]: from sklearn.metrics import jaccard_similarity_score
```

```

def get_jaccard_similarity(bag_of_words_matrix):
    """
    Get jaccard similarities for neighboring documents

    Parameters
    -----
    bag_of_words : 2-d Numpy Nddarray of int
        Bag of words sentiment for each document
        The first dimension is the document.
        The second dimension is the word.

    Returns
    -----
    jaccard_similarities : list of float
        Jaccard similarities for neighboring documents
    """

    # TODO: Implement
    bag_of_words = bag_of_words_matrix.astype(bool)
    jaccard_similarities = [jaccard_similarity_score(u,v) for u, v in zip(bag_of_words, bag_of_words)]

    return jaccard_similarities

```

```
project_tests.test_get_jaccard_similarity(get_jaccard_similarity)
```

Using the get_jaccard_similarity function, let's plot the similarities over time.

```

In [ ]: # Get dates for the universe
file_dates = {
    ticker: [ten_k['file_date'] for ten_k in ten_ks]
    for ticker, ten_ks in ten_ks_by_ticker.items()}

jaccard_similarities = {
    ticker: {
        sentiment_name: get_jaccard_similarity(sentiment_values)
        for sentiment_name, sentiment_values in ten_k_sentiments.items()}
    for ticker, ten_k_sentiments in sentiment_bow_ten_ks.items()}

project_helper.plot_similarities(
    [jaccard_similarities[example_ticker][sentiment] for sentiment in sentiments],
    file_dates[example_ticker][1:],
    'Jaccard Similarities for {} Sentiment'.format(example_ticker),
    sentiments)

```

1.5.4 TFIDF

using the sentiment word lists, let's generate sentiment TFIDF from the 10-k documents. Implement `get_tfidf` to generate TFIDF from each document, using sentiment words as the terms. You can ignore words that are not in `sentiment_words`.

```
In [ ]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
def get_tfidf(sentiment_words, docs):
    """
    Generate TFIDF values from documents for a certain sentiment

    Parameters
    -----
    sentiment_words: Pandas Series
        Words that signify a certain sentiment
    docs : list of str
        List of documents used to generate bag of words

    Returns
    -----
    tfidf : 2-d Numpy Nddarray of float
        TFIDF sentiment for each document
        The first dimension is the document.
        The second dimension is the word.
    """

    # TODO: Implement
    vectorizer = TfidfVectorizer(vocabulary = sentiment_words)
    tfidf = vectorizer.fit_transform(docs).toarray()

    return tfidf
```

```
project_tests.test_get_tfidf(get_tfidf)
```

Using the `get_tfidf` function, let's generate the TFIDF values for all the documents.

```
In [ ]: sentiment_tfidf_ten_ks = {}
```

```
for ticker, ten_ks in ten_ks_by_ticker.items():
    lemma_docs = [' '.join(ten_k['file_lemma']) for ten_k in ten_ks]

    sentiment_tfidf_ten_ks[ticker] = {
        sentiment: get_tfidf(sentiment_df[sentiment_df[sentiment]]['word'], lemma_docs)
        for sentiment in sentiments}
```

```
project_helper.print_ten_k_data([sentiment_tfidf_ten_ks[example_ticker]], sentiments)
```


1.5.5 Cosine Similarity

Using the TFIDF values, we'll calculate the cosine similarity and plot it over time. Implement `get_cosine_similarity` to return the cosine similarities between each tick in time. Since the input, `tfidf_matrix`, is a TFIDF vector for each time period in order, you just need to compute the cosine similarities for each neighboring vector.

```
In [ ]: from sklearn.metrics.pairwise import cosine_similarity
```

```
def get_cosine_similarity(tfidf_matrix):
    """
    Get cosine similarities for each neighboring TFIDF vector/document

    Parameters
    -----
    tfidf : 2-d Numpy Nddarray of float
            TFIDF sentiment for each document
            The first dimension is the document.
            The second dimension is the word.

    Returns
    -----
    cosine_similarities : list of float
                        Cosine similarities for neighboring documents
    """

    # TODO: Implement
    cos_similarity= cosine_similarity(tfidf_matrix[0:], tfidf_matrix[1:])

    return cos_similarity[0].tolist()
```

```
project_tests.test_get_cosine_similarity(get_cosine_similarity)
```

Let's plot the cosine similarities over time.

```
In [ ]: cosine_similarities = {
    ticker: {
        sentiment_name: get_cosine_similarity(sentiment_values)
        for sentiment_name, sentiment_values in ten_k_sentiments.items()}
    for ticker, ten_k_sentiments in sentiment_tfidf_ten_ks.items()}

project_helper.plot_similarities(
    [cosine_similarities[example_ticker][sentiment] for sentiment in sentiments],
    file_dates[example_ticker][1:],
    'Cosine Similarities for {} Sentiment'.format(example_ticker),
    sentiments)
```

1.6 Evaluate Alpha Factors

Just like we did in project 4, let's evaluate the alpha factors. For this section, we'll just be looking at the cosine similarities, but it can be applied to the jaccard similarities as well. `### Price Data`
Let's get yearly pricing to run the factor against, since 10-Ks are produced annually.

```
In [ ]: pricing = pd.read_csv('../data/project_5_yr/yr-quotemedia.csv', parse_dates=['date'])
pricing = pricing.pivot(index='date', columns='ticker', values='adj_close')

pricing
```

1.6.1 Dict to DataFrame

The `alphalens` library uses dataframes, so we we'll need to turn our dictionary into a dataframe.

```
In [ ]: cosine_similarities_df_dict = {'date': [], 'ticker': [], 'sentiment': [], 'value': []}

for ticker, ten_k_sentiments in cosine_similarities.items():
    for sentiment_name, sentiment_values in ten_k_sentiments.items():
        for sentiment_value, sentiment_value in enumerate(sentiment_values):
            cosine_similarities_df_dict['ticker'].append(ticker)
            cosine_similarities_df_dict['sentiment'].append(sentiment_name)
            cosine_similarities_df_dict['value'].append(sentiment_value)
            cosine_similarities_df_dict['date'].append(file_dates[ticker][1:][sentiment_value])

cosine_similarities_df = pd.DataFrame(cosine_similarities_df_dict)
cosine_similarities_df['date'] = pd.DatetimeIndex(cosine_similarities_df['date']).year
cosine_similarities_df['date'] = pd.to_datetime(cosine_similarities_df['date'], format='%Y')

cosine_similarities_df.head()
```

1.6.2 Alphalens Format

In order to use a lot of the `alphalens` functions, we need to aligned the indices and convert the time to unix timestamp. In this next cell, we'll do just that.

```
In [ ]: import alphalens as al

factor_data = {}
skipped_sentiments = []

for sentiment in sentiments:
    cs_df = cosine_similarities_df[(cosine_similarities_df['sentiment'] == sentiment)]
    cs_df = cs_df.pivot(index='date', columns='ticker', values='value')
```

```

try:
    data = al.utils.get_clean_factor_and_forward_returns(cs_df.stack(), pricing, qua
    factor_data[sentiment] = data
except:
    skipped_sentiments.append(sentiment)

if skipped_sentiments:
    print('\nSkipped the following sentiments:\n{}'.format('\n'.join(skipped_sentiments))
factor_data[sentiments[0]].head()

```

1.6.3 Alphalens Format with Unix Time

Alphalens's `factor_rank_autocorrelation` and `mean_return_by_quantile` functions require unix timestamps to work, so we'll also create factor dataframes with unix time.

```

In [ ]: unixt_factor_data = {
    factor: data.set_index(pd.MultiIndex.from_tuples(
        [(x.timestamp(), y) for x, y in data.index.values],
        names=['date', 'asset']))
    for factor, data in factor_data.items()}

```

1.6.4 Factor Returns

Let's view the factor returns over time. We should be seeing it generally move up and to the right.

```

In [ ]: ls_factor_returns = pd.DataFrame()

for factor_name, data in factor_data.items():
    ls_factor_returns[factor_name] = al.performance.factor_returns(data).iloc[:, 0]

(1 + ls_factor_returns).cumprod().plot()

```

1.6.5 Basis Points Per Day per Quantile

It is not enough to look just at the factor weighted return. A good alpha is also monotonic in quantiles. Let's look at the basis points for the factor returns.

```

In [ ]: qr_factor_returns = pd.DataFrame()

for factor_name, data in unixt_factor_data.items():
    qr_factor_returns[factor_name] = al.performance.mean_return_by_quantile(data)[0].iloc[:, 0]

(10000*qr_factor_returns).plot.bar(
    subplots=True,
    sharey=True,
    layout=(5,3),
    figsize=(14, 14),
    legend=False)

```

1.6.6 Turnover Analysis

Without doing a full and formal backtest, we can analyze how stable the alphas are over time. Stability in this sense means that from period to period, the alpha ranks do not change much. Since trading is costly, we always prefer, all other things being equal, that the ranks do not change significantly per period. We can measure this with the **Factor Rank Autocorrelation (FRA)**.

```
In [ ]: ls_FRA = pd.DataFrame()

        for factor, data in unixt_factor_data.items():
            ls_FRA[factor] = al.performance.factor_rank_autocorrelation(data)

        ls_FRA.plot(title="Factor Rank Autocorrelation")
```

1.6.7 Sharpe Ratio of the Alphas

The last analysis we'll do on the factors will be sharpe ratio. Let's see what the sharpe ratio for the factors are. Generally, a Sharpe Ratio of near 1.0 or higher is an acceptable single alpha for this universe.

```
In [ ]: daily_annualization_factor = np.sqrt(252)

        (daily_annualization_factor * ls_factor_returns.mean() / ls_factor_returns.std()).round(
```

That's it! You've successfully done sentiment analysis on 10-ks! ## Submission Now that you're done with the project, it's time to submit it. Click the submit button in the bottom right. One of our reviewers will give you feedback on your project with a pass or not passed grade. You can continue to the next section while you wait for feedback.